

```
"""
```

プログラム 1 : CNNモデルを構築し、画像データを学習させます。

```
"""
```

```
%tensorflow_version 1.x

import os
import matplotlib.pyplot as plt
import keras
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D, BatchNormalization
from keras import regularizers
from keras.callbacks import ReduceLROnPlateau, ModelCheckpoint
from keras.preprocessing.image import ImageDataGenerator
import random

os.chdir('/content/drive/My Drive/Python/CNN/')

num_classes = 10
im_rows = 32
im_cols = 32
in_shape = (im_rows, im_cols, 3)

# データを読み込む
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# データを正規化
x_train = x_train.astype("float32") / 255
x_test = x_test.astype("float32") / 255
# ラベルデータをOne-Hot形式に変換
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
# 検証データとテストデータに分ける
mask = list(range(10000))
random.shuffle(mask)
x_val = x_test[mask[:5000]]
y_val = y_test[mask[:5000]]
```

```

y_val = y_test[mask[:5000]]
x_test = x_test[mask[5000:]]
y_test = y_test[mask[5000:]]
# データ拡張
train_datagen = ImageDataGenerator(
    rotation_range=30,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.1,
    fill_mode="nearest",
    horizontal_flip=True
)
train_generator = train_datagen.flow(x_train, y_train, batch_size=512)

# モデルを定義
model = Sequential()
model.add(Conv2D(128, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05),
    input_shape = in_shape))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Conv2D(128, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Dropout(0.3))

model.add(Conv2D(256, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Conv2D(256, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Conv2D(256, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size = (2, 2)))

```

```

model.add(Dropout(0.3))

model.add(Conv2D(512, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Conv2D(512, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Conv2D(512, (3, 3), padding="same", kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Dropout(0.3))

```

```

model.add(Flatten())
model.add(Dense(512, kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("relu"))
model.add(Dropout(0.3))
model.add(Dense(num_classes, kernel_regularizer= regularizers.l1(1e-05)))
model.add(BatchNormalization())
model.add(Activation("softmax"))
model.summary()

```

モデルをコンパイル

```

model.compile(
    loss = "categorical_crossentropy",
    optimizer = "adadelta",
    metrics = ["accuracy"])

```

学習を実行

```

reduce_lr = ReduceLROnPlateau(monitor='accuracy', factor=0.8, patience=5, ¥
    verbose=1, mode='auto', epsilon=0.01, cooldown=0, min_lr=0.0001)
filepath = 'weights_max.hdf5'
checkpoint = ModelCheckpoint(filepath, monitor='val_accuracy', verbose=1, ¥
    save_best_only=True, mode='auto')
hist = model.fit_generator(train_generator,
    steps_per_epoch=300

```

```
steps_per_epoch=1000,  
epochs=100,  
verbose=1,  
validation_data=(x_val, y_val),  
callbacks=[reduce_lr, checkpoint])
```

モデルを評価

```
score = model.evaluate(x_test, y_test, verbose = 1)  
print("正確率=", score[1], "loss=", score[0])
```

学習の様子をグラフへ描画

```
plt.plot(hist.history["accuracy"])  
plt.plot(hist.history["val_accuracy"])  
plt.title("Accuracy")  
plt.legend(["train", "test"], loc = "upper left")  
plt.show()
```

```
plt.plot(hist.history["loss"])  
plt.plot(hist.history["val_loss"])  
plt.title("Loss")  
plt.legend(["train", "test"], loc = "upper left")  
plt.show()
```

学習結果を保存

```
model.save("cifar10-cnn-weight.h5")
```



TensorFlow 1.x selected.
Using TensorFlow backend.
Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 14s 0us/step
WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/resource_variable_ops.py:1630: calling BaseResourceVariable._Instructions for updating:
If using Keras pass *_constraint arguments to layers.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4070: The name tf.nn.max_pool is deprecated

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 32, 32, 128)	3584
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 128)	512
activation_1 (Activation)	(None, 32, 32, 128)	0
conv2d_2 (Conv2D)	(None, 32, 32, 128)	147584
batch_normalization_2 (Batch Normalization)	(None, 32, 32, 128)	512
activation_2 (Activation)	(None, 32, 32, 128)	0
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 128)	0
dropout_1 (Dropout)	(None, 16, 16, 128)	0
conv2d_3 (Conv2D)	(None, 16, 16, 256)	295168
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 256)	1024
activation_3 (Activation)	(None, 16, 16, 256)	0
conv2d_4 (Conv2D)	(None, 16, 16, 256)	590080
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 256)	1024
activation_4 (Activation)	(None, 16, 16, 256)	0
conv2d_5 (Conv2D)	(None, 16, 16, 256)	590080

Layer	Kernel	Stride	Output	Params
batch_normalization_5	(Batch Normalization)	(None, 16, 16, 256)	1024	
activation_5	(Activation)	(None, 16, 16, 256)	0	
max_pooling2d_2	(MaxPooling2D)	(None, 8, 8, 256)	0	
dropout_2	(Dropout)	(None, 8, 8, 256)	0	
conv2d_6	(Conv2D)	(None, 8, 8, 512)	1180160	
batch_normalization_6	(Batch Normalization)	(None, 8, 8, 512)	2048	
activation_6	(Activation)	(None, 8, 8, 512)	0	
conv2d_7	(Conv2D)	(None, 8, 8, 512)	2359808	
batch_normalization_7	(Batch Normalization)	(None, 8, 8, 512)	2048	
activation_7	(Activation)	(None, 8, 8, 512)	0	
conv2d_8	(Conv2D)	(None, 8, 8, 512)	2359808	
batch_normalization_8	(Batch Normalization)	(None, 8, 8, 512)	2048	
activation_8	(Activation)	(None, 8, 8, 512)	0	
max_pooling2d_3	(MaxPooling2D)	(None, 4, 4, 512)	0	
dropout_3	(Dropout)	(None, 4, 4, 512)	0	
flatten_1	(Flatten)	(None, 8192)	0	
dense_1	(Dense)	(None, 512)	4194816	
batch_normalization_9	(Batch Normalization)	(None, 512)	2048	
activation_9	(Activation)	(None, 512)	0	
dropout_4	(Dropout)	(None, 512)	0	
dense_2	(Dense)	(None, 10)	5130	

batch_normalization_10 (Batch Normalization)	(None, 10)	40
activation_10 (Activation)	(None, 10)	0

=====
 Total params: 11,738,546
 Trainable params: 11,732,382
 Non-trainable params: 6,164

/usr/local/lib/python3.6/dist-packages/keras/callbacks/callbacks.py:998: UserWarning: `epsilon` argument is deprecated and will be removed, use
 warnings.warn("`epsilon` argument is deprecated and
 WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:422: The name tf.global_variables is deprecated

Epoch 1/100
 300/300 [=====] - 184s 613ms/step - loss: 2.8823 - accuracy: 0.5348 - val_loss: 3.7329 - val_accuracy: 0.2384

Epoch 00001: val_accuracy improved from -inf to 0.23840, saving model to weights_max.hdf5

Epoch 2/100
 300/300 [=====] - 172s 572ms/step - loss: 2.1503 - accuracy: 0.7140 - val_loss: 2.1939 - val_accuracy: 0.6610

Epoch 00002: val_accuracy improved from 0.23840 to 0.66100, saving model to weights_max.hdf5

Epoch 3/100
 300/300 [=====] - 172s 573ms/step - loss: 1.7511 - accuracy: 0.7733 - val_loss: 1.5888 - val_accuracy: 0.7966

Epoch 00003: val_accuracy improved from 0.66100 to 0.79660, saving model to weights_max.hdf5

Epoch 4/100
 300/300 [=====] - 172s 574ms/step - loss: 1.4855 - accuracy: 0.8026 - val_loss: 1.4691 - val_accuracy: 0.7894

Epoch 00004: val_accuracy did not improve from 0.79660

Epoch 5/100
 300/300 [=====] - 173s 575ms/step - loss: 1.2938 - accuracy: 0.8230 - val_loss: 1.3414 - val_accuracy: 0.7960

Epoch 00005: val_accuracy did not improve from 0.79660

Epoch 6/100
 300/300 [=====] - 173s 577ms/step - loss: 1.1546 - accuracy: 0.8396 - val_loss: 1.2223 - val_accuracy: 0.8048

Epoch 00006: val_accuracy improved from 0.79660 to 0.80480, saving model to weights_max.hdf5

Epoch 7/100
 300/300 [=====] - 172s 574ms/step - loss: 1.0665 - accuracy: 0.8493 - val_loss: 1.0964 - val_accuracy: 0.8330

Epoch 00007: val_accuracy improved from 0.80480 to 0.83300, saving model to weights_max.hdf5

Epoch 8/100

Epoch 00096: val_accuracy did not improve from 0.94060

Epoch 97/100

300/300 [=====] - 171s 570ms/step - loss: 0.2322 - accuracy: 0.9938 - val_loss: 0.4669 - val_accuracy: 0.9362

Epoch 00097: val_accuracy did not improve from 0.94060

Epoch 98/100

300/300 [=====] - 171s 570ms/step - loss: 0.2309 - accuracy: 0.9938 - val_loss: 0.4880 - val_accuracy: 0.9342

Epoch 00098: val_accuracy did not improve from 0.94060

Epoch 99/100

300/300 [=====] - 171s 571ms/step - loss: 0.2284 - accuracy: 0.9942 - val_loss: 0.4799 - val_accuracy: 0.9324

Epoch 00099: ReduceLROnPlateau reducing learning rate to 0.04398046135902405.

Epoch 00099: val_accuracy did not improve from 0.94060


Epoch 100/100

300/300 [=====] - 171s 571ms/step - loss: 0.2258 - accuracy: 0.9944 - val_loss: 0.4950 - val_accuracy: 0.9308

Epoch 00100: val_accuracy did not improve from 0.94060

5000/5000 [=====] - 3s 651us/step

正確率= 0.925599992275238 loss= 0.5341672481536865




```
"""
```

プログラム2：プログラム2で学習済みのCNNモデルを使って、画像を判定させます。

```
"""
```

```
import os
import cv2
import numpy as np
from keras.models import load_model
import matplotlib.pyplot as plt
labels = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
im_shape = (32, 32, 3)
```

```
os.chdir('/content/drive/My Drive/Python/CNN/')
```

```
# モデルデータを読み込み
```

```
model = load_model("cifar10-cnn-weight.h5")
```

```
# OpenCVを使って画像を読み込み
```

```
im = cv2.imread("test-car.jpg")
```

```
# 色空間を変換して、リサイズ
```

```
im = cv2.cvtColor(im, cv2.COLOR_BGR2RGB)
```

```
im = cv2.resize(im, (32, 32))
```

```
plt.imshow(im)
```

```
plt.show()
```

```
# MLPで学習した画像データに合わせる
```

```
im = im.reshape(im_shape).astype("float32") / 255
```

```
# 予測する
```

```
r = model.predict(np.array([im]), batch_size = 32, verbose = 1)
```

```
res = r[0]
```

```
# 結果を表示する
```

```
for i, acc in enumerate(res):
```

```
    print(labels[i], "=", int(acc * 100))
```

```
print("----")
```

```
print("予測した結果:", labels[res.argmax()])
```





1/1 [=====] - 1s 736ms/step

airplane = 0
automobile = 99
bird = 0
cat = 0
deer = 0
dog = 0
frog = 0
horse = 0
ship = 0
truck = 0

予測した結果: automobile